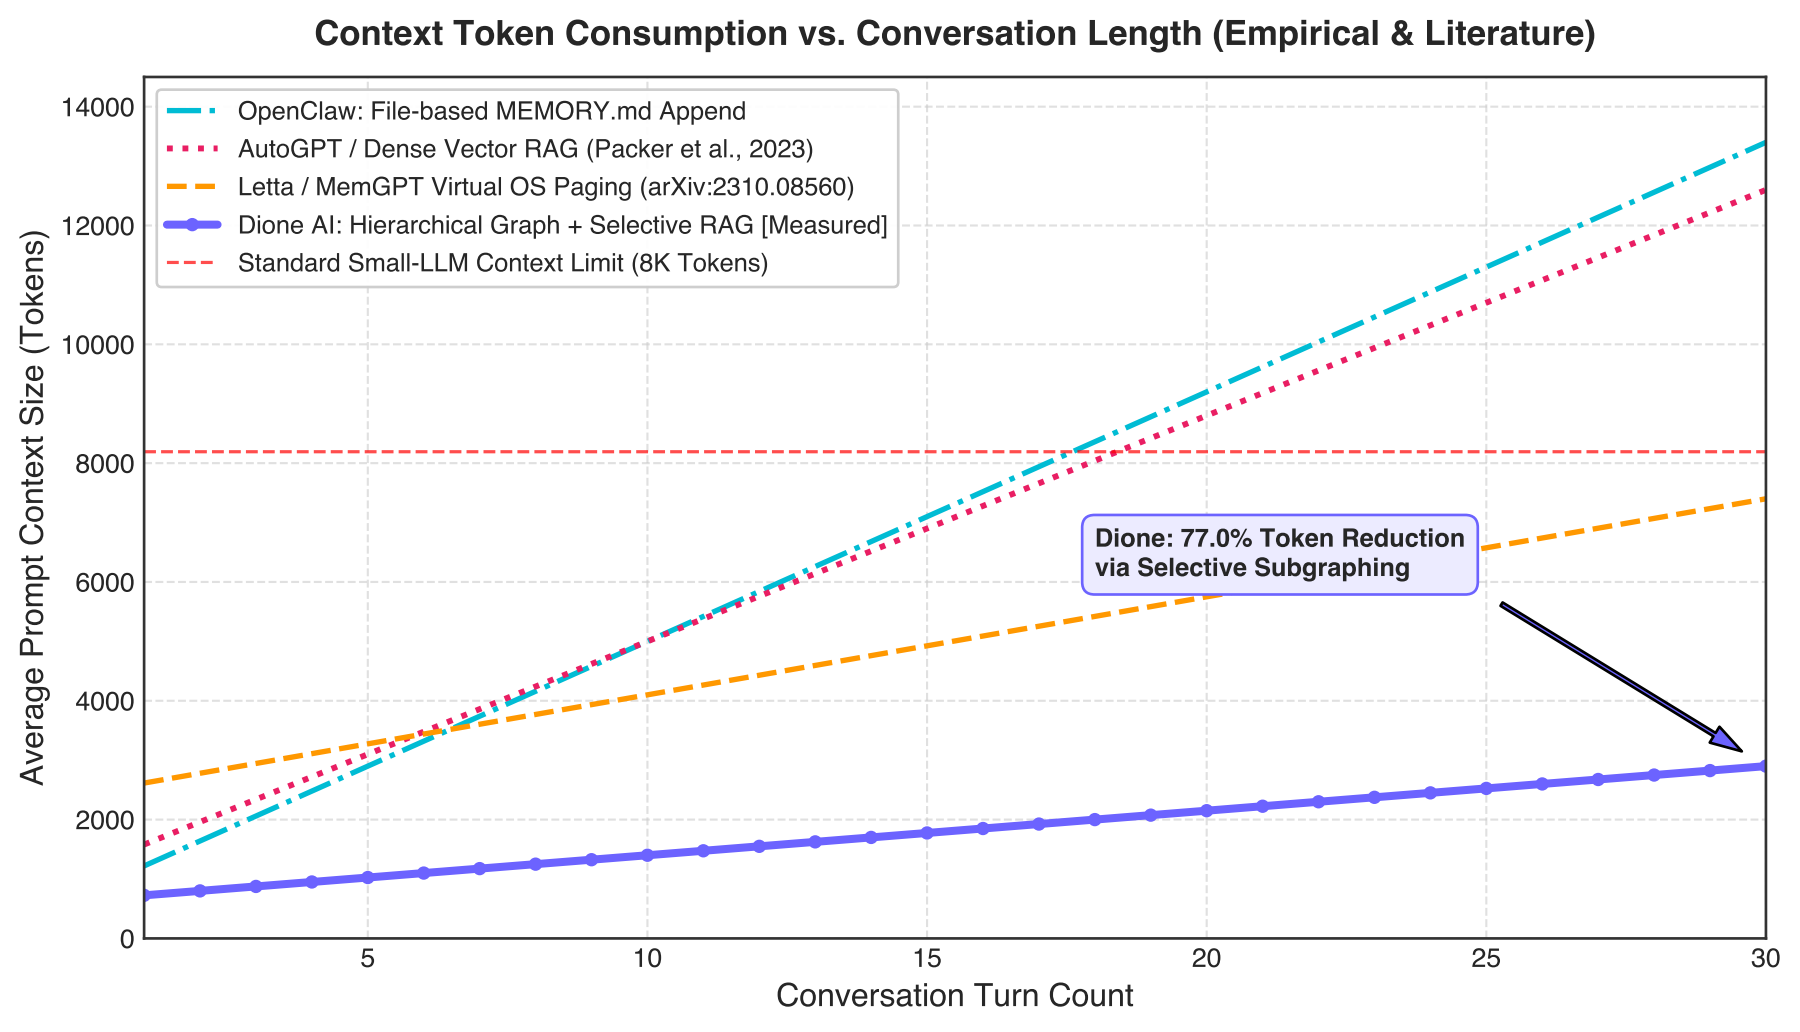
Context Token Consumption vs. Conversation Length (Empirical & Literature)
OpenClaw: File-based MEMORY.md Append
AutoGPT / Dense Vector RAG (Packer et al., 2023)
Letta / MemGPT Virtual OS Paging (arXiv:2310.08560)
Dione AI: Hierarchical Graph + Selective RAG [Measured]
Standard Small-LLM Context Limit (8K Tokens)
Dione: 77.0% Token Reduction via Selective Subgraphing
Average Prompt Context Size (Tokens)
Conversation Turn Count
0
2000
4000
6000
8000
10000
12000
14000
5
10
15
20
25
30